


Fig. 2 — GMedMamba: spectral pathway conditioning a 4-stage spatial backbone (gmedmamba.py)



Verified by running gmedmamba.py: 27.42 M params — gmedmamba_tiny (RGB, 6 classes, dims 96/192/384/768). 2.77 M — gmedmamba_hsi_small (patch_size=1, any C incl. 826 bands, 3 classes): C never appears in a parameter shape.
Fusion type is selectable per stage (film default / gated / cross_attention / multiplicative / residual); dynamic band selection and stage-wise spectral → spatial refinement are both on by default.
Scope: this figure covers gmedmamba.py only (GMedMamba/GMedMambaBackbone + ClassificationHead). Reconstruction decoding and EMA weight averaging are separate modules attached by the training scripts, not part of this file.