


Fig. 3 — GMedMamba-R: the hierarchical backbone replaced by one weight-shared recursive core (gmedmamba.py)



Verified by running gmedmamba.py (gmedmamba_trm(), default trm_mixer="ss2d"): 0.531 M params vs 27.42 M hierarchical tiny — 51.6× fewer. mlp mixer: 0.443 M (61.9×); attention: 0.373 M (73.5×).

Core f is called (n_latent+1)×n_improve×segments = 7×3×4 = 84 times per forward pass; each call runs trm_core_layers=2 RecursiveMambaBlocks (mixer + GatedMLP, post-RMSNorm residuals).

Scope: this figure covers gmedmamba.py only (GMedMambaRecursive/GMedMambaRecursiveBackbone + RecursiveHead). EMA weight averaging (trm_ema_rate is only a config field here) is applied by the external trainer, not by this file.